CS 466/666: Algorithm Design and Analysis	University of Waterloo: Fall 2025

Lecture 8

September 30, 2025

Instructor: Sepehr Assadi

Disclaimer: *These notes have not been subjected to the usual scrutiny reserved for formal publications. They may be distributed outside this class only with the permission of the Instructor.*

Topics of this Lecture

1 Introduction to Linear Programming

1.1 Optimization Problems

Optimization problems are everywhere: How to design the best schedule of classes? How to go from point A to point B in the quickest possible way? How to transfer information in a network at the highest rate possible? How to fit the largest number of packages in a delivery truck? These, and numerous examples alike, are all optimization problems.

For our purpose, we can model optimization problems formally as follows.

> **Definition 1.** Let Ω be a **domain** and $f : \Omega \to \mathbb{R}$ be an **objective function** that assign value to the elements of this domain. Moreover, let $C \subseteq \Omega$ denote a subset of the domain identified by the **constraints** of the problem. The goal is now to solve the following problem:
>
> $$\max_{x \in \Omega} \quad f(x) \qquad \text{subject to} \qquad x \in C.$$

In words, Definition 1 simply defines optimization problems as finding an element of the domain Ω that also belongs to the "feasible" subset C and maximizes the value of $f(x)$. Notice that it is also possible for an optimization problem to involve *minimizing* the objective function instead; however, since minimizing $f(x)$ and maximizing $-f(x)$ are equivalent, there is no need for a further definition.

Attempting to solve optimization problems at this level of generality is virtually hopeless. For instance,

- Consider the following family of optimization problems $\{\Omega_n, f_n : \Omega_n \to \mathbb{R}, C_n\}_{n\in\mathbb{N}}$: for any integer $n \in \mathbb{N}$, we take Ω_n to be the set of all n-states Turing machines, C_n to be the subset of n-states Turing machines that eventually halt on the input 0, and $f_n(x)$ for each Turing machine $x \in \Omega_n$ be the function that counts the number of 1's output by x on input 0 (we set $f(x) = 0$ if x outputs any symbol other than 1).

  Now, maximizing $f_n(x)$ subject to $x \in C_n$ is equivalent to the infamous *busy beaver* problem which is a well-known undecidable problem. Thus, it is impossible to design an algorithm for solving this optimization problem in any finite time.

- Another, perhaps "scarier" example, is the following optimization problem on only four variables:

$$\min_{x=(x_1,x_2,x_3,x_4)\in\mathbb{N}^4} \ |x_1^{x_4} + x_2^{x_4} - x_3^{x_4}| \qquad \text{subject to} \qquad x_1, x_2, x_3 \geqslant 1 \text{ and } x_4 \geqslant 3.$$

  Checking whether the optimal solution to this optimization problem has value 0 or not is equivalent to proving *Fermat's Last Theorem*!

Given these, it is clear that we need to impose certain restriction on our optimization problems before we can hope to study them further.

1.2 Linear Optimization (a.k.a. Linear Programming)

In this course, we will examine a canonical family of optimization problems: *Linear Programs (LPs)*[1]. In a linear program, the domain Ω is $\mathbb{R}^n$, the objective function is always a linear function $c^T \cdot x$ (or equivalently $\langle c, x \rangle$), and set C is identified by a series of linear constraints of the form $\langle a, x \rangle \geqslant b$. Formally,

> **Definition 2.** In a linear program (LP), we have an n-dimensional vector of **variables** $x \in \mathbb{R}^n$, an n-dimensional **objective function** $c \in \mathbb{R}^n$, a set m **constraints** given via a $m \times n$ matrix $A \in \mathbb{R}^{m\times n}$ and an m-dimensional vector $b \in \mathbb{R}^m$. The goal is to solve the following problem:
> $$\max_{x\in\mathbb{R}^n} \ c^\top \cdot x \qquad \text{subject to} \qquad A \cdot x \geqslant b.$$

A simple example of a linear program is the following:

$$\max_{(x_1,x_2,x_3)\in\mathbb{R}^3} \ x_1 + 2x_2 + 3x_3$$
$$\text{subject to} \quad x_1 + x_2 + x_3 \leqslant 3$$
$$x_2 + x_3 \leqslant 2$$
$$x_1, x_2, x_3 \geqslant 0.$$

The optimal solution to this LP is $(x_1, x_2, x_3) = (1, 0, 2)$ which achieves the value of 7 (see if you can *verify* this is indeed an optimal solution).

There are many applications of linear programs:

(a) We can *solve* numerous "day-to-day life" problems, ranging from simple scheduling to complicated scientific computing tasks, by modeling them via LPs and solving them using available LP solvers.

(b) We can use LPs as a tool toward *designing* "generic" algorithms for various problems, for instance, in *combinatorial optimization.*

(c) We can use LPs as an *analytical* tool for proving theorems (particularly in combinatorics) by exploiting strong properties of LPs such as duality theorems (that we will learn about in the next lectures).

---

[1]The term 'programming' here has nothing to do with computer programming and instead is used similar to 'scheduling'.

2   Applications of Linear Programming

Let us make the above applications more concrete via several examples. We will see applications (a) and (b) in this lecture and postpone application (c) to the next one.

2.1   An Application of Type (a): Creating a "Balanced" Diet

The following is a textbook example of applications of LPs in solving "day-to-day life" problems. Suppose we are interested in creating a balanced diet in terms of only its contents in vitamins A,B, C, its fiber, and its protein. We are further provided by the dietary guidelines on how much of each of these our meal should consists of, as well as a list of foods together with their vitamins/fiber/protein contents. On top of this, we are given the price for each of these foods. This is basically a table of the following form:

	Required amount	Food 1 (e.g., carrot)	Food 2 (e.g., pizza)	$\cdots$	Food n
Vitamin A	r_A	a_1	a_2		a_n
Vitamin B	r_B	b_1	b_2		b_n
Vitamin C	r_C	c_1	c_2		c_n
Fiber	r_F	f_1	f_2		f_n
Protein	r_P	p_1	p_2		p_n
Price	$-$	$price_1$	$price_2$		$price_n$

Table 1: This table was supposed to be filled with actual numbers but this course is not about dietary guidelines and I had no idea what numbers to put here; so, we instead write them as the above parameters.

The goal is now to find a diet as a combination of these foods that satisfy all the required amounts of vitamins, fiber, and protein, and has the minimum price possible.

We can write this as an LP as follows:

- For $i \in [n]$, let x_i be a variable denoting how much of food i we will include;

- For $i \in [n]$, let a_i, b_i, c_i, f_i, p_i, denote vitamins A,B,C, fiber, and protein content of food i, respectively;

- For $i \in [n]$, let $price_i$ denote the price we have to pay for each unit of food i;

- Finally, let r_A, r_B, r_C, r_F, r_P denote the required amount of vitamins A,B,C, fiber, and protein.

Note that the last three lines are *not* variables, rather, numbers that are (supposedly) in Table 1

We can now write the following LP for solving our problem.

$$\min_{x=(x_1,\ldots,x_n)\in\mathbb{R}^n} \quad \sum_{i=1}^{n} price_i \cdot x_i$$

subject to

$$\sum_{i=1}^{n} a_i \cdot x_i \geqslant r_A \quad \sum_{i=1}^{n} b_i \cdot x_i \geqslant r_B, \quad \sum_{i=1}^{n} c_i \cdot x_i \geqslant r_C, \quad \sum_{i=1}^{n} f_i \cdot x_i \geqslant r_F, \quad \sum_{i=1}^{n} p_i \cdot x_i \geqslant r_P$$

$$x_i \geqslant 0 \qquad \forall i \in [n].$$

(We again emphasize that in the above LP, the only variables are $x_1, \ldots, x_n$; remaining parameters are all numbers that we can fill up using Table 1 assuming we had the actual numbers.)

We will leave verifying this LP actually solves the original problem to the reader as an easy exercise. At this point, we can simply give this LP to a standard LP solver and obtain the solution to our problem.

2.2 Another Application of Type (a): Linear Regression

In the **linear regression** problem, we have a set of m points from the n-dimensional space $\mathbb{R}^n$:

$$(a_{1,1}, a_{1,2}, \ldots, a_{1,n}), (a_{2,1}, a_{2,2}, \ldots, a_{2,n}) \cdots, (a_{m,1}, a_{m,2}, \ldots, a_{m,n}).$$

We can denote these points by a matrix $A \in \mathbb{R}^{m \times n}$. The points are associated with real numbers in $\mathbb{R}$:

$$(b_1, b_2, \ldots, b_m).$$

Similarly, we can denote these numbers by a vector $b \in \mathbb{R}^m$.

Think of each column of this matrix as a "feature" or an "attribute" of the points, and think of each real number b_i as the "label" of the point $a_i := (a_{i,1}, \ldots, a_{i,n})$. Our goal is to find an "explanation" of these labels in terms of the feature using a *linear* function. Formally, we would like to find a *hyperplane* $x \in \mathbb{R}^n$ such that given any point a_i, we can recover the label of a_i, namely, b_i, via $\langle a_i , x \rangle$.

Of course, such a hyperplane can only exist if the system of linear equations $A \cdot x = b$ has a solution. However, this is not generally guaranteed (especially because we often have $m \gg n$). Thus, the goal is to find a hyperplane x with minimal "distance" from b. This leads to the family of regression problems wherein the goal is to solve the following optimization problem:

$$\min_{x \in \mathbb{R}^n} \|A \cdot x - b\|,$$

where we can pick the $\|\cdot\|$ differently depending on the application. Two popular choices of the norm are ℓ_2-norm and ℓ_1-norm. You might have seen ℓ_2-regression elsewhere which admits a closed form solution

$$x = (A^\top A)^{-1} A^\top b;$$

this requires $A^\top A$ to be non-singular which is often the case (if not, we have to take the pseudo-inverse here instead). Regardless, this question is out of the scope of this lecture. We will instead focus on ℓ_1-regression that can be solved via linear programming. Writing this problem as a LP is not entirely trivial so we are going to do it step by step.

Step 1. Recall that the ℓ_1-norm of a vector $v \in \mathbb{R}^m$ is defined as:

$$\|v\|_1 := \sum_{i=1}^{n} |v_i|.$$

Thus, the ℓ_1-regression problem is the following optimization problem:

$$\min_{x \in \mathbb{R}^n} \sum_{i=1}^{m} |\langle a_i , x \rangle - b_i|.$$

This is of course yet not a linear program because the objective function is not a linear function. But we made the problem somewhat simpler without changing it at all (it is straightforward to check the problems are equivalent so far).

Step 2. We are going to define m new variables $z_1, \ldots, z_m$, one per each row of the matrix A. We write our optimization problem now as:

$$\min_{x \in \mathbb{R}^n, z \in \mathbb{R}^m} \quad \sum_{i=1}^{m} z_i$$
$$\text{subject to} \quad z_i = \max\left(\langle a_i , x \rangle - b_i, b_i - \langle a_i , x \rangle\right) \quad \forall i \in [m].$$

This is still equivalent to the previous problem because for every $i \in [m]$,

$$|\langle a_i , x \rangle - b_i| = \max\left(\langle a_i , x \rangle - b_i, b_i - \langle a_i , x \rangle\right),$$

by definition; enforcing z_i's being equal to this ensures that the problem remains the same. We are now one step closer since our objective function now is linear although we still have many non-linear constraints.

Step 3. The only (slightly) non-trivial step is this one where we are going to relax the constraints a bit without violating feasibility as follows:

$$\min_{x \in \mathbb{R}^n, z \in \mathbb{R}^m} \quad \sum_{i=1}^{m} z_i$$
$$\text{subject to} \quad z_i \geqslant \max\left(\langle a_i , x \rangle - b_i, b_i - \langle a_i , x \rangle\right) \quad \forall i \in [m].$$

It may not be entirely clear that this problem is the same as above, in particular because we actually *expanded* the feasible region. However, we still have the following claim.

Claim 3. *In any* optimal *solution of this problem, we have that for every* $i \in [m]$,

$$z_i = \max\left(\langle a_i , x \rangle - b_i, b_i - \langle a_i , x \rangle\right);$$

Proof. Suppose some z_i is strictly larger; then reduce it by a tiny $\varepsilon > 0$ which does not violate the constraint but reduces the objective function, contradicting the optimality. $\square$

This implies that optimal solutions of the problems in steps 2 and 3 coincide and thus we can still focus on solving this problem. We are still not done because our constraints are not yet linear.

Step 4. Finally, we can "linearize" the constraints simply as follows:

$$\min_{x \in \mathbb{R}^n, z \in \mathbb{R}^m} \quad \sum_{i=1}^{m} z_i$$
$$\text{subject to} \quad z_i \geqslant \langle a_i , x \rangle - b_i \quad \text{and} \quad z_i \geqslant b_i - \langle a_i , x \rangle \quad \forall i \in [m].$$

This is now truly a linear program and is clearly equivalent to the previous step because for all $i \in [m]$:

$$z_i \geqslant \max\left(\langle a_i , x \rangle - b_i, b_i - \langle a_i , x \rangle\right) \iff z_i \geqslant \langle a_i , x \rangle - b_i \quad \text{and} \quad z_i \geqslant b_i - \langle a_i , x \rangle$$

Wrap-up. Using the above four steps, we can write any ℓ_1-regression problem with m points in n dimension as a LP with $n + m$ variables and $2m$ constraints. Solving this LP then gives us the optimal solution to the ℓ_1-regression problem as well.

2.3   An Application of Type (b): The Bipartite Matching Problem

We consider yet another textbook example of LPs, this time for solving algorithmic problems. In the **bipartite matching** problem, we are given a bipartite graph $G = (L \sqcup R, E)$ with bipartition L and R of vertices. A **matching** M in G is any collection of edges that do not share any vertices (hence, they 'match' vertices in L to vertices in R so that each participating vertex in the matching is matched to precisely one other vertex). In the bipartite matching problem, the goal is to find a maximum matching, namely, a one with the largest number of edges.

Let us define an *integral* 0/1 variable $y_e \in \{0, 1\}$ for each edge $e \in E$ (we emphasize that these variables are *not* in $\mathbb{R}^n$, but rather we forced them to be only 0 or 1). Think of $y_e = 1$ as showing that we pick the edge e in our maximum matching and $y_e = 0$ as we do not.

To ensure that the edges we pick form a matching, we need to make sure we do not pick more than one edge per vertex. We can write this as the following constraint:

$$\forall u \in L : \quad \sum_{e \ni u} y_e \leqslant 1 \qquad \text{and} \qquad \forall v \in R : \quad \sum_{e \ni v} y_e \leqslant 1.$$

Finally, our task of maximizing the number of edges in the matching becomes

$$\max \sum_{e \in E} y_e.$$

Putting these together, we get the following optimization problem:

$$\max_{y} \quad \sum_{e \in E} y_e$$

$$\text{subject to} \quad \forall u \in L : \quad \sum_{e \ni u} y_e \leqslant 1$$

$$\forall v \in R : \quad \sum_{e \ni v} y_e \leqslant 1$$

$$\forall e \in E : \quad y_e \in \{0, 1\}\,.$$

The above optimization problem is *not* an LP (even though it may very much look like it). The problem is with the very last constraint of $y_e \in \{0, 1\}$ that forces the program to only use integer values for the variables. Such a program is instead called an **Integer Linear Program (ILP)**. Unfortunately, unlike LPs, ILPs are much harder to solve or reason about. For one thing, solving general ILPs is an NP-hard problem (you can prove this easily; try!).

There is a general recipe in these situations: what if we *relax* the integrality constraint to obtain the following LP instead?

$$\max_{x \in \mathbb{R}^E} \quad \sum_{e \in E} x_e$$

$$\text{subject to} \quad \forall u \in L : \quad \sum_{e \ni u} x_e \leqslant 1$$

$$\forall v \in R : \quad \sum_{e \ni v} x_e \leqslant 1$$

$$\forall e \in E : \quad 0 \leqslant x_e \leqslant 1.$$

Such an LP is typically called the **LP relaxation** of the original ILP (or directly the original problem). But, now it is no longer clear that solving this LP can tell us much about the original problem. Of course, the optimal value of this LP is always an upper bound on the optimal value of the original ILP simply because any solution to that ILP is a feasible solution here as well. But can it be that this relaxation made the problem so different that these values now differ drastically from each other?

The answer to this question *for general LP relaxations and ILPs* can be *Yes*; we will see examples of this soon in the course. However, we are now going to prove that for the specific case of the bipartite matching problem, this is not the case.

Proposition 4. *Any feasible fractional solution $x \in \mathbb{R}^E$ to the LP for bipartite matching can be* rounded *to a feasible integral solution $y \in \{0, 1\}^E$ without decreasing the objective value.*

Before getting to the proof, notice a direct corollary of this proposition. We can now solve the bipartite matching LP relaxation to find an optimal solution x. We know that the objective value of x is at least as large as the optimal solution of the ILP. By Proposition 4, we can then round x to an integral solution y of the same value, thus implying that y is also optimal for the ILP. This then gives us a maximum matching of the original graph and solves the problem.

Proof of Proposition 4. Fix the feasible solution $x \in \mathbb{R}^E$. If x is already integral, we are done by taking $y = x$. Otherwise, perform the following step first.

Step one: cycle canceling. Find a cycle in the support of x, denoted by $\text{supp}(x) := \{e \in E : x_e > 0\}$, that contains a fractional value. If no such cycle is found go to the next step, otherwise let C be the cycle and let $e \in C$ be the edge with the smallest fractional value x_e in the cycle. Let $\delta = x_e$. Now update x as follows. Let $C = e_1, e_2, e_3, \ldots, e_\ell$ denote the edges of the cycle ordered such that $e_1 = e$ is the picked edge with $x_e = \delta$. Update

$$x_{e_1} \leftarrow x_{e_1} - \delta, \quad x_{e_2} \leftarrow x_{e_2} + \delta, \quad x_{e_3} \leftarrow x_{e_3} - \delta, \ldots,$$

Figure 1: An illustration of cycle canceling.

namely, alternately decrease and increase x-values of edges by δ. See Figure 1 for an illustration.

Since we are in a bipartite graph, the length of C must be even (recall that bipartite graphs are precisely those that have no odd cycles). This ensures that (i) the value of $\sum_{e \ni w} x_e$ remains the same for every vertex $w \in L \cup R$ in the graph since we increased one edge of w and decreased another one by the same value; and (ii) value of $\sum_{e \in E} x_e$ remains the same by the same reasoning. Moreover, support of x is now at least one edge smaller because the original edge e we started with now has $x_e = 0$.

Thus, through this transformation, we only *shrink* the support of x without decreasing the objective value. As the support is finite, this process has to end eventually and at that point, we can no longer have any cycles inside it.

Step two: rounding forests. Now that there are no cycles left in the support of x, we have that $\mathrm{supp}(x)$ is a forest (namely, a graph without cycles). Any part of this support that is integral has to be disjoint from the rest because integral parts can only form a matching and thus vertices have degree exactly one (they form a collection of vertex-disjoint edges).

Let F be the fractional part of the support. Pick any arbitrary leaf-nodes in the forest F (namely, a vertex of degree one). Let v be this leaf-node u be its parent, and $e = (u, v)$ be the edge between them. Let $x_e = 1$ and reduce the x-value of all other edges incident on u to become 0. Firstly, since v was a leaf-node and thus had no other edges, this transformation does not make x infeasible. Secondly, the total value of x incident on vertex u previously was at most 1 (to ensure feasibility of the original x) and by this change it has become exactly 1. Thus, we did not reduce the value of x at all. Finally, all edges incident on u have integral values and thus are removed from F. We can continue this process and at every step we obtain a feasible solution of the same value or higher and eventually transform the entire F into an integral solution.

Let y be this final solution which is now integral and has value as large as x. This concludes the proof. $\square$

> **Remark.** In general, rounding a fractional solution to an integral one does not necessarily need to be "lossless". Finding the best rounding scheme (even the best LP relaxation of a problem) is one of the main topics studied in *approximation algorithms*, and we will revisit this later in the course. The gap between the optimal ILP value and its LP relaxation is referred to as the **integrality gap**.

2.4   Another Application of Type (b): The Bipartite Vertex Cover Problem

Let $G = (L \sqcup R, E)$ be a bipartite graph. A vertex cover in G is a set of vertices S that *cover* all the edges, i.e., any edge has at least one endpoint in S. Let us write this LP relaxation for bipartite vertex cover:

$$\min_{y \in \mathbb{R}^V} \quad \sum y_v$$

7

$$\text{subject to} \quad y_u + y_v \geqslant 1 \qquad \forall\, (u, v) \in E,$$
$$y_v \geqslant 0 \qquad\qquad \forall\, v \in L \sqcup R.$$

As before, *if* assume the values in this LP are integral, then we get a program for minimum vertex cover, because $y_v = 1$ can be interpreted as picking the vertex v in the solution, and the $y_u + y_v \geqslant 1$ constraints ensures that from every edge we are picking at least one vertex. Of course, in general, there is no reason this LP picks its solution integrally and thus the LP more accurately corresponds to the *fractional* vertex cover problem. However, we show that the optimum value of this LP is equal to the minimum vertex cover of G.

Proposition 5. *Any feasible solution $y \in \mathbb{R}^V$ of the LP for bipartite vertex cover can be rounded to a vertex cover of the input graph with size at most $\sum_{v \in V} y_v$.*

Proof. Pick $\theta \in (0, 1)$ uniformly at random and define the sets:

$$L^* = \{u \mid u \in L \text{ and } y_u \geqslant \theta\} \qquad \text{and} \qquad R^* = \{v \mid v \in R \text{ and } y_v \geqslant 1 - \theta\}.$$

We claim the set $L^* \cup R^*$ is a vertex cover of G because for an arbitrary edge $(u, v) \in E$ such that $u \in L, v \in R$, at least one of $u \in L^*$ or $v \in R^*$ is true. If $u \notin L^*$ and $v \notin R^*$, then we have $y_u < \theta$ and $y_v < 1 - \theta$, thus $y_u + y_v < 1$ which contradicts the feasibility of y for the LP.

We now bound the size of the solution in expectation.

$$\mathbb{E}\,|L^* \cup R^*| = \sum_{u \in L} \Pr(u \in L^*) + \sum_{v \in R} \Pr(v \in R^*)$$
$$= \sum_{u \in L} \Pr(y_u \geqslant \theta) + \sum_{v \in R} \Pr(y_v \geqslant 1 - \theta) \qquad \text{(by the definition of } L^* \text{ and } R^*\text{)}$$
$$= \sum_{u \in L} y_u + \sum_{v \in R} y_v \qquad\qquad\qquad (\theta \text{ is chosen uniformly at random})$$
$$= \sum_{v \in V} y_v.$$

This concludes the proof. $\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad\quad$ $\square$

Remark. An important remark is in order here. While the above approach gives a randomized algorithm for rounding the value of the LP, a slightly more careful view of the same analysis implies that we can actually replace the random choice of $\theta \in (0, 1)$ with an *arbitrary* choice instead (for instance, simply take $\theta = 1/2$).

This is because in this algorithm, we always output a vertex cover. Our analysis shows that the expected size of this vertex cover is at most that of the minimum vertex cover. But then this implies that for *every* choice of θ, the size of the returned vertex cover should be equal to the minimum vertex cover size; otherwise, to "balance" the expectation, some choice of θ should result in a value which is even less than the minimum vertex cover size which is not possible[a].

---
[a]If for a random variable Y, $\mathbb{E}[Y] = \min_{y \in \mathrm{supp}(Y)} y$, then we should have that $\Pr(Y = \min_{y \in \mathrm{supp}(Y)} y) = 1$.