

Lecture 24

December 2, 2025

Instructor: Sepehr Assadi

Disclaimer: These notes have not been subjected to the usual scrutiny reserved for formal publications. They may be distributed outside this class only with the permission of the Instructor.

Topics of this Lecture

1	Background Tools from the Last Lecture	1
2	An $\tilde{O}(m\sqrt{n})$ Time Algorithm for SSSP	2
2.1	Proof of Lemma 5: The Scaling Lemma	3
3	The Final $\tilde{O}(m)$ Time Algorithm	5

We finish this course by going over the algorithm of [BNW22] for the negative weight shortest path. In particular, we will see a proof of the following theorem.

Theorem 1 ([BNW22]). *There is a randomized algorithm that for any graph $G = (V, E, w)$ with no negative cycle, finds single-source shortest paths from a given vertex $s \in V$ in $O(m \log^4 n \cdot \log(nW))$ expected time where $W := \max_e |w(e)|$, namely, the largest absolute value of any edge weight.*

1 Background Tools from the Last Lecture

Let us recap the following two main lemmas from the last lecture.

Lemma 2. *There is an algorithm for solving SSSP in any graph in $O(m \log n \cdot \bar{b}_s)$ time, where \bar{b}_s is the average negative-hop distance of all vertices from s (check Lecture 23 for the definitions).*

Lemma 3 (Directed Low Diameter Decomposition [BNW22]). *There is a randomized algorithm that given any directed graph $G = (V, E, w)$ with non-negative integer weights and an integer $D \geq 1$, outputs a set of edges E_{rem} with the following properties:*

- Let C be any strongly connected component (SCC) of $G \setminus E_{\text{rem}}$. Then, C has a “weak diameter” at most D :

$$\forall u, v \in C \quad \text{dist}_G(u, v) \leq D \quad \text{and} \quad \text{dist}_G(v, u) \leq D.$$

- For any edge $e \in E$,

$$\Pr(e \in E_{\text{rem}}) = \frac{O(\log^2 n)}{D} + n^{-10}.$$

The algorithm runs in $O(m \log^3 n)$ time deterministically (in fact, the runtime is $O(m \log^2 n + n \log^3 n)$ but the distinction is not important for us in this lecture).

2 An $\tilde{O}(m\sqrt{n})$ Time Algorithm for SSSP

We start by proving a weaker version of [Theorem 1](#) which has many of the key ideas. We then sketch how we can extend this algorithm to prove [Theorem 1](#) also. In particular, our goal now is to prove the following simpler theorem which gives an $\tilde{O}(m\sqrt{n})$ runtime¹ for SSSP.

Theorem 4 (Weaker version of [Theorem 1](#) [BNW22]). *There is a randomized algorithm that for any graph $G = (V, E, w)$ with no negative cycle, finds single-source shortest paths from a given vertex $s \in V$ in $O(m\sqrt{n} \log^2 n \cdot \log(nW))$ expected time.*

The general framework of the proof is as follows: suppose we start with G and a weight function w such that $w(e) \geq -2B$ for all $e \in E$ and some integer $B \geq 1$. We will find a price function ϕ such that after applying it, we will have $w_\phi(e) \geq -B$ for all $e \in E$; i.e., the most negative-weight edge is now at most half as negative as before. We then show that repeating this step for $O(\log W)$ iterations is enough to obtain weights that can be thought of as essentially non-negative. This approach is often called **scaling** and is a classical technique in algorithm design. Let us now formalize this further.

Lemma 5 (Scaling Lemma). *There is a randomized algorithm that given any graph $G = (V, E, w)$ where $w(e) \geq -2B$ for every edge $e \in E$, outputs a price function function ϕ such that $w_\phi(e) \geq -B$ for every edge $e \in E$. The algorithm has expected $O(m\sqrt{n} \log^2 n)$ time.*

Let us see how to use this lemma now to obtain an $\tilde{O}(m\sqrt{n})$ time algorithm for SSSP.

Proof of [Theorem 4](#) using [Lemma 5](#). Given $G = (V, E, w)$ with $W := \max_e |w(e)|$, we first update the weight function w so that $w(e) \leftarrow n \cdot w(e)$. Clearly, this does not change the shortest path structures (nor change positivity/negativity of any edge). We then run [Lemma 5](#) repeatedly for $t := \log(nW)$ iterations on the graph G to obtain price functions $\phi_1, \phi_2, \dots, \phi_t$, where

$$w(e) \geq -nW \implies w_{\phi_1}(e) \geq -\frac{nW}{2} \implies w_{\phi_2}(e) \geq -\frac{nW}{2^2} \implies \dots \implies w_{\phi_t} \geq -\frac{nW}{2^t} = -1,$$

where each ‘ \implies ’ corresponds to running the algorithm of [Lemma 5](#) once. Note that here, each price function ϕ_i is applied on top of the price function ϕ_{i-1} , i.e., is obtained by adding the price function of [Lemma 5](#) to the price function ϕ_i .

At this point, we define a new weight function w' wherein $w'(e) = w_{\phi_t}(e) + 1$. In principle, this can potentially change the shortest path structure. Nevertheless, we prove in our special case, this cannot happen. Suppose P and Q are two s - v paths in G (under the original weight function w) and that $w(P) < w(Q)$. We argue that $w'(P) < w'(Q)$ also. We have,

$$w'(P) = w_{\phi_t}(P) + |P| = w(P) + \phi_t(s) - \phi_t(v) + |P| \leq w(P) + (n-1) + \phi_t(s) - \phi_t(v),$$

since P can have $n-1$ edges at most. On the other hand

$$w'(Q) = w_{\phi_t}(Q) + |Q| = w(Q) + \phi_t(s) - \phi_t(v) + |Q| \geq w(Q) + \phi_t(s) - \phi_t(v).$$

But recall that since we updated the weights w by multiplying them by n , if $w(P) < w(Q)$, then in fact, $w(P) \leq w(Q) - n$ even. Thus, we continue to have $w'(P) < w'(Q)$ also as desired.

Since w' is a non-negative weight function, we can simply run Dijkstra’s algorithm on G, w' and solve SSSP in $O(m \log n)$ time at this point (and by the previous argument and correctness of price functions, we get the solution is correct). Thus, the runtime of the algorithm is $O(m\sqrt{n} \cdot \log^2 n \cdot \log(nW))$ as desired. \square

¹Recall that $\tilde{O}(f) := O(f \cdot \text{poly log}(f))$.

2.1 Proof of Lemma 5: The Scaling Lemma

We update the graph by adding a vertex s^* that is connected to every vertex with an edge of weight $-B$. Throughout, we use the same set of vertices $\{s^*\} \cup V$ but with different subset of edges (subsets of E which is now updated to include (s^*, v) -edges as well) and different weight functions. In particular, define the following two additional weight functions:

$$w^{2B} : E \rightarrow \mathbb{Z} \quad \text{wherein } w^{2B}(e) = w(e) + 2B \text{ for all edges } e \in E;$$

$$w^B : E \rightarrow \mathbb{Z} \quad \text{wherein } w^B(e) = w(e) + B \text{ for all edges } e \in E.$$

Note that both these weight functions can potentially destroy the shortest path structure (but that will not be a concern for us because we will only use them to compute a price function). Moreover w^{2B} is now a non-negative weight function. The algorithm, at a high level, is as follows.

Algorithm 1. The high-level description of the algorithm of Lemma 5. The parameter d below will be set later. The steps of the algorithm will be explained in more detail later.

1. Compute a LDD of $G^{2B} = (\{s^*\} \cup V, E, w^{2B})$ with diameter $D = dB$ using Lemma 3. Let C_1, \dots, C_k be the SCCs and E_{rem} be the removed edges of the LDD.
2. Use the weights w^B (and not w^{2B}) to find a price function ϕ_1 that makes all edges inside C_i 's non-negative in $w_{\phi_1}^B$.
3. Use the updated weights $w_{\phi_1}^B$ to find a price function ϕ_2 that additionally makes the DAG edges of $G \setminus E_{\text{rem}}$ non-negative in $w_{\phi_2}^B$.
4. Use the updated weights $w_{\phi_2}^B$ to find a price function ϕ_3 that additionally makes the edges in E_{rem} non-negative in $w_{\phi_3}^B$.
5. Return w_{ϕ_3} (and not $w_{\phi_3}^B$) as a weight function that satisfy $w_{\phi_3}(e) \geq -B$ for all $e \in E$.

We will now go over different steps of this algorithm in detail.

Step 1: LDD computation. Recall that an LDD is only defined for graphs with non-negative weights. Since $w(e) \geq -2B$ by assumption and $w^{2B}(e) = w(e) + 2B$ by definition, we have w^{2B} is a non-negative weight function. As such, it is valid to apply Lemma 3 and obtain a set E_{rem} of edges such that any SCC C of $G \setminus E_{\text{rem}}$ satisfies:

$$\forall u, v \in C \quad \text{dist}_{G^{2B}}(u, v) \leq dB \quad \text{and} \quad \text{dist}_{G^{2B}}(v, u) \leq dB, \quad (1)$$

and for any edge in G

$$\Pr(e \in E_{\text{rem}}) = \frac{O(\log^2 n)}{dB} \cdot w^{2B}(e). \quad (2)$$

This step takes $O(m \log^3 n)$ time.

Step 2: Fixing SCC edges. Consider the graph $G_1 = (\{s^*\} \cup V, E_1, w^B)$ with weight function w^B where $E_1 \subseteq E$ only contains the edges between SCCs of $G \setminus E_{\text{rem}}$ (i.e., remaining edges after removing E_{rem} and DAG edges). Note that the edges of s^* to all other vertices have weight 0 under w^B (as they had weight $-B$ under w). We claim that the shortest paths from s^* in this graph have “few” negative edges.

Claim 6. For any $v \in V$, the hop distance of s^* to v in G_1 is less than d .

Proof. Consider a shortest path P_{s^*v} in G_1 from s^* to v in G_1 . We know that $w^B(P_{s^*v}) \leq 0$ as s^* is connected to v by an edge of weight 0. If $w^B(P_{s^*v}) = 0$, the hop distance of s^* to v will simply be one by taking the (s^*, v) edge directly. Otherwise, we have $w^B(P_{s^*v}) < 0$ and thus P_{s^*v} starts by going from s^* to some vertex u in the same SCC as v and then taking the shortest path P_{uv} inside this SCC (recall that the only edges of G_1 are the ones inside SCCs). We will argue that P_{uv} can have $< d$ edges.

Suppose towards a contradiction that P_{uv} contains at least d edges in G_1 . Then, under the original weight function w , we have,

$$w(P_{uv}) = w^B(P_{uv}) - |P_{uv}| \cdot B < 0 - dB = -dB,$$

since $|P_{uv}| \geq d$ by our assumption. On the other hand, since u and v are both inside the same SCC of $G \setminus E_{\text{rem}}$, by Eq (3), we have

$$\text{dist}_{G^{2B}}(v, u) \leq dB.$$

This implies that there exists some path Q_{vu} in G (and not necessarily G_1) such that

$$w(Q_{vu}) = w^{2B}(v, u) - |Q_{vu}| \cdot 2B \leq dB - 1.$$

Putting these two implies that in the original graph G , we can go from u to v with a path of weight $< -dB$ and from v to u with a path of weight $< dB$. This implies that we can start from u and return to it by paying a total weight $< -dB + dB < 0$, implying that there must be a negative cycle in G . But this is a contradiction with the statement of Theorem 4 that implied there is no negative cycle in G . \square

Combining Claim 6 with Lemma 2 implies that we can find s^* -shortest paths in G_1 in $O(m \log n \cdot d)$ time. We will then define, for any $v \in V$,

$$\phi_1(v) = \text{dist}_{G_1}(s^*, v).$$

This implies that for any edge $(u, v) \in G_1$,

$$w_{\phi_1}^B(u, v) = w^B(u, v) + \phi_1(u) - \phi_1(v) = w^B(u, v) + \text{dist}_{G_1}(s^*, u) - \text{dist}_{G_1}(s^*, v) \geq 0,$$

where the last inequality is by triangle inequality since (u, v) is an edge of G_1 . Thus, we made all SCC edges non-negative under $w_{\phi_1}^B$.

Step 3: Fixing DAG edges. This step is quite straightforward: we compute a topological ordering of the SCCs of the graph $G \setminus E_{\text{rem}}$ in $O(m + n)$ and denote them by C_1, \dots, C_k . For any $v \in C_i$, we define

$$\phi_2(v) = \phi_1(v) + (k - i) \cdot \max_{e \in E} |w_{\phi_1}^B(e)|.$$

Note that for any edge (u, v) inside the same cluster C_i , we have

$$w_{\phi_2}^B(u, v) = w^B(u, v) + \phi_2(u) - \phi_2(v) = w^B(u, v) + \phi_1(u) - \phi_1(v) = w_{\phi_1}^B(u, v) \geq 0,$$

as we proved in the last part. For any edge $u \in C_i$ and $v \in C_j$ for $j > i$,

$$w_{\phi_2}^B(u, v) = w^B(u, v) + \phi_2(u) - \phi_2(v) = w^B(u, v) + \phi_1(u) - \phi_1(v) + \max_{e \in E} |w_{\phi_1}^B(e)| = w_{\phi_1}^B(u, v) + \max_{e \in E} |w_{\phi_1}^B(e)| \geq 0.$$

Finally, since we are working with a topological ordering of a DAG there are no other edges, and thus all edges, except for E_{rem} , have become non-negative in $w_{\phi_2}^B$.

Step 4: Fixing E_{rem} edges. We now consider s^* -shortest paths in the entire graph G under the updated weight function $w_{\phi_2}^B$. We claim that these shortest paths also contain only a “few” negative edges on average.

Claim 7. For any $v \in V$, the expected negative hop distance of s^* to v in G under the weight function $w_{\phi_2}^B$ is less than $\frac{O(h_{s^*}(v) \cdot \log^2 n)}{d}$, where $h_s(v)$ is the hop distance of s^* to v under the weight function w^B .

Proof. As in [Claim 6](#), we consider a path P_{s^*v} that goes from s^* to some vertex u and then take P_{uv} with $w_{\phi_2}^B(P_{uv}) < 0$ (otherwise, the hop distance of s^* to v will be one). Note that P_{uv} is also the shortest path from u to v in w^B itself also since price functions do change the shortest path structure. But under w^B , we could have again gone from s^* to v with a weight of 0, and thus $w^B(P_{uv}) < 0$. This implies that

$$w^{2B}(P_{uv}) = w^B(P_{uv}) + |P_{uv}| \cdot B \leq h_{s^*}(v) \cdot B.$$

At the same time, the number of negative edges in P_{uv} under $w_{\phi_2}^B$ is at most equal to $|P_{uv} \cap E_{\text{rem}}|$ as previous steps made sure all other edges are non-negative. Thus,

$$\begin{aligned} \mathbb{E} [\text{negative hop distance of } s^* \text{ to } v \text{ in } w_{\phi_2}^B] &\leq \mathbb{E} |P_{uv} \cap E_{\text{rem}}| \\ &= \sum_{e \in P_{uv}} \Pr(e \in E_{\text{rem}}) && (\text{by linearity of expectation}) \\ &= \sum_{e \in P_{uv}} \frac{O(\log^2 n)}{dB} w^{2B}(e) && (\text{by Eq (4)}) \\ &= \frac{O(\log^2 n)}{dB} \cdot h_{s^*}(v) \cdot B && (\text{by the above bound on } w^{2B}(P_{uv})) \\ &= \frac{O(h_{s^*}(v) \cdot \log^2 n)}{d}, \end{aligned}$$

as desired. \square

Since the hop distances are always at most $n - 1$, by combining [Claim 7](#) and [Lemma 2](#), we can find s^* -shortest path in the entire G under the weight function $w_{\phi_2}^B$ in $O(m \log^3 n \cdot \frac{n}{d})$ expected time. By setting

$$\phi_3(v) = \phi_2(v) + \text{dist}_{w_{\phi_2}^B}(s^*, v),$$

for all $v \in V$, we can make all edges of G non-negative under the weight function $w_{\phi_3}^B$. Finally, this implies that under the original weight function w but with the price function ϕ_3 , for any $e \in E$, we have

$$w_{\phi_3}(e) = w_{\phi_3}^B(e) - B \geq -B.$$

The expected runtime of the algorithm is now

$$O(m \log^3 n + m \log n \cdot d + m \log^3 n \cdot \frac{n}{d}),$$

and thus by setting $d = \sqrt{n} \log n$, we obtain the expected runtime of

$$O(m \sqrt{n} \log^2 n),$$

concluding the proof of [Lemma 5](#).

3 The Final $\tilde{O}(m)$ Time Algorithm

The algorithm in [Theorem 1](#) can also be obtained in a very similar manner, using the following improved scaling lemma.

Lemma 8 (Improved Scaling Lemma). *There is a randomized algorithm that given any graph $G = (V, E, w)$ where $w(e) \geq -2B$ for every edge $e \in E$, outputs a price function function ϕ such that $w_\phi(e) \geq -B$ for every edge $e \in E$. The algorithm has expected $O(m \log^4 n)$ time.*

[Theorem 1](#) follows from [Lemma 8](#) the same exact way [Theorem 4](#) followed from [Lemma 5](#). We now show how to prove [Lemma 8](#).

The idea behind the proof of [Lemma 8](#) is to introduce one more level of recursion: instead of balancing the time took in Step 2 and 4 of [Algorithm 1](#) that led to an $\tilde{O}(m\sqrt{n})$ time, we will make Step 4 much faster and then recurse on the graphs of Step 2. The improvement obtained in Step 2 comes from another metric: the negative hop distances of the SCCs still drop by a factor of two, and thus the recursion depth will only be $O(\log n)$ which makes our algorithm fast enough.

More formally, the algorithm is as follows. Note that we again use the weight functions w^{2B} and w^B and also add the vertex s^* to the graph as before.

Algorithm 2. The high-level description of the algorithm of [Lemma 8](#). The input is a graph $G = (\{s^*\} \cup V, E, w)$ with an additional parameter Δ promised to be an upper bound on the hop distances between all pairs of reachable vertices in V (ignoring s^*) under the weight function w^B . The algorithm returns a price function ϕ such that $w_\phi^B(e) \geq 0$ for all $e \in E$. The steps of the algorithm are also explained in more detail later.

1. If $\Delta \leq 1$, run a base case algorithm (explained below) and return.
2. Compute a LDD of $G^{2B} = (\{s^*\} \cup V, E, w^{2B})$ with diameter $D = d \cdot B$ using [Lemma 3](#) for a parameter $d = \Delta/2$ to be fixed explicitly later. Let C_1, \dots, C_k be the SCCs and E_{rem} be the removed edges of the LDD.
3. Use the weights w^B (and not w^{2B}) and recurse on the graphs $G_i = (\{s^*\} \cup C_i, E[s^* \cup C_i], w^B)$ with parameter $\Delta/2$ to find a price function ϕ_1 that makes edges inside C_i 's non-negative in $w_{\phi_1}^B$.
4. Use the updated weights $w_{\phi_1}^B$ to find a price function ϕ_2 that additionally makes the DAG edges of $G \setminus E_{\text{rem}}$ non-negative in $w_{\phi_2}^B$.
5. Use the updated weights $w_{\phi_2}^B$ to find a price function ϕ_3 that additionally makes the edges in E_{rem} non-negative in $w_{\phi_3}^B$.
6. Return $w_{\phi_3}^B$ as a weight function that satisfy $w_{\phi_3}^B(e) \geq 0$ for all $e \in E$.

Step 1: Base case. For the base case, we simply need to run [Lemma 2](#) to find s^* -shortest paths and set

$$\phi(v) = \text{dist}_{G^B}(s^*, v).$$

The correctness follows as before as these distances make w^B non-negative and thus $w(e) \geq -B$ for all $e \in E$. Moreover, the runtime is only $O(m \log n)$ by [Lemma 2](#) and the promise that the shortest path between every pair of vertices inside V uses at most one hop. (Technically speaking, we could have just run one iteration of the Bellman-Ford algorithm and solve the problem in $O(m)$ time but the difference is inconsequential).

Step 2: LDD computation. As before, w^{2B} is non-negative and so it is valid to apply [Lemma 3](#) and obtain a set E_{rem} of edges such that any SCC C of $G \setminus E_{\text{rem}}$ satisfies:

$$\forall u, v \in C \quad \text{dist}_{G^{2B}}(u, v) \leq dB \quad \text{and} \quad \text{dist}_{G^{2B}}(v, u) \leq dB, \tag{3}$$

and for any edge in G

$$\Pr(e \in E_{\text{rem}}) = \frac{O(\log^2 n)}{dB} \cdot w^{2B}(e) = \frac{O(\log^2 n)}{\Delta \cdot B} \cdot w^{2B}(e). \tag{4}$$

This step takes $O(m \log^3 n)$ time.

Step 3: Fixing SCC edges. We do exactly as in [Algorithm 1](#) and by [Claim 6](#), have that under w^B , the negative hop diameter of every C_i will be $d = \Delta/2$. This means that the recursive call in this step is run with a correct parameter and thus by induction, we will find a price function ϕ_1 that ensures $w_{\phi_1}^B(e) \geq 0$ for all e in the SCCs.

Step 4: Fixing DAG edges. This step is exactly as before and can be done in $O(m + n)$ time.

Step 5: Fixing E_{rem} edges. Again, we exactly as in [Algorithm 1](#) and by [Claim 7](#), have that under $w_{\phi_2}^B$, the negative hop distance of any vertex in expectation is

$$O(\log^2 n \cdot \frac{h_{s^*}(v)}{d}) = O(\log^2 n) \cdot \frac{\Delta}{\Delta/2} = O(\log^2 n),$$

using the fact that under w^B (by our initial assumption in the recursion), hop diameter of the graph is Δ and since we set $d = \Delta/2$. This means in that this step can now be implemented in $O(m \log^3 n)$ expected time using [Lemma 2](#).

In conclusion, the algorithm correctly finds a price function ϕ such that w_ϕ^B is non-negative and thus for every $e \in E$, $w_\phi(e) \geq -B$ as desired.

For the runtime analysis, the algorithm reduces the value of Δ by a factor of two each time, and we would be calling it with $\Delta = n - 1$ at the beginning since any pair of reachable vertices can have at most $n - 1$ hops between them. This means there are $O(\log n)$ level of recursion. Each level also takes $O(m \log^3 n)$ expected time at most, leading to a total of $O(m \log^4 n)$ expected time. This concludes the proof of [Lemma 8](#) and the entire proof of [Theorem 1](#).

References

- [BNW22] Aaron Bernstein, Danupon Nanongkai, and Christian Wulff-Nilsen. Negative-weight single-source shortest paths in near-linear time. In *2022 IEEE 63rd annual symposium on foundations of computer science (FOCS)*, pages 600–611. IEEE, 2022. [1](#), [2](#)